

Punggol Digital District Digital Twin

Optional Case Study in Tool Integration, Agentic Workflows, and Safe Operations

Tang Liang

DigiPen Institute of Technology Singapore

AI Application Development Bootcamp – Optional Reading

How to Use This Reading

Documented foundation

PDD, the Open Digital Platform (ODP), its digital twin, integrated building systems, sensors, simulation, and testbed role are publicly documented.

Teaching extension

The MCP interfaces, agent controllers, robot workflows, tool names, and incident scenarios in this deck are conceptual designs—not claims about deployed PDD or SIT systems.

Purpose

Study how a real district platform can become the environment for tool-integrated and agentic applications, while keeping operational control safe.

Reading Roadmap

- 1 **PDD and the Open Digital Platform**
- 2 **From Digital Twin to Application Architecture**
- 3 **Conceptual Agentic Extension**
- 4 **Three Service Scenarios**
- 5 **Safety, Governance, and Reflection**

PDD at a Glance

A mixed-use digital district

Punggol Digital District integrates a business park with the Singapore Institute of Technology campus and community amenities.

Scale and ambition

JTC describes a 50-hectare district designed for about 28,000 jobs and 12,000 students, supported by shared digital infrastructure.

Source: JTC, [Punggol Digital District](#).


Why a District-Scale Digital Twin?

- Buildings, transport, energy, security, facilities, and public spaces affect one another.
- Individual dashboards produce silos; a shared operational view exposes cross-system relationships.
- Historical playback supports diagnosis; simulation supports planning before physical action.
- A common platform can let public agencies, operators, researchers, and solution providers test services.

Engineering shift

The twin is not merely a 3D model. It is a continuously updated operational representation connected to data and systems.

ODP: The District's Digital Backbone



Key idea

ODP provides an interoperability and abstraction layer so applications need not integrate separately with every building subsystem.

Source: JTC, [Open Digital Platform: the digital backbone of PDD](#).

What Can Be Integrated?

Domain	Example data	Example capability
Facilities	equipment status, alarms, work orders	monitor assets and coordinate maintenance
Energy	cooling, electricity, PV, storage, demand	optimise resource use and compare scenarios
Mobility	car parks, traffic, autonomous systems	coordinate access and movement
Security	access-control and security events	correlate events across systems
Environment	temperature, weather, occupancy	understand comfort and operating conditions

Source: JTC, [ODP press release](#), 2018.

The Twin Supports More Than Monitoring

Observe

Live 3D operational view and sensor/system status.

Explain

Historical playback and evidence for diagnosing an event.

Explore


Simulation and virtual testing before changes reach the physical district.

Do not confuse the layers

A digital twin represents and simulates the world; an agent decides how to pursue a goal; deterministic controllers execute protected operations.

2. From Digital Twin to Application Architecture

A Shared World Model



Shared world model

Applications reason from consistent identifiers, locations, relationships, status, and time—instead of inventing their own incompatible view of the district.

Control Plane vs Data Plane

Control plane

Commands, queries, approvals, tickets, policy decisions, and audit records. Usually lower volume and semantically meaningful.

Data plane

Video, LiDAR, robot telemetry, building telemetry, and high-frequency sensor streams. Often high volume and latency-sensitive.

Design rule

Do not send every video frame or motor-control message through an LLM or MCP. Give the agent summaries and bounded high-level capabilities.

Which Integration Mechanism Fits?

Mechanism	Strong fit	Limitation
Direct API/SDK	stable production service, performance, precise contract	different integration for each API
CLI	local tools, scripts, developer automation	text parsing and host permissions need care
Event bus	telemetry and asynchronous events at scale	not itself an LLM tool interface
MCP	discoverable, typed tools/resources for AI clients	extra layer; not ideal for all data paths
ROS / field protocols	robot and equipment communication	domain-specific; not a reasoning layer

Hybrid answer

A district application uses several mechanisms. MCP is one adapter at the AI-facing control boundary, not the universal network.

Where MCP Could Fit—Conceptual Extension



- MCP can describe tools and resources to the AI application.
- The application service translates them into existing APIs, workflows, or commands.
- Identity, policy, validation, and logging remain outside the LLM.

Status

This is a teaching architecture, not a statement that PDD uses MCP.

Unity or Web 3D: View, Simulation, Interaction

- A 3D client can render rooms, assets, robots, warnings, routes, and historical events.
- The backend/twin should remain the operational source of truth.
- The rendering client should call controlled services rather than directly modifying live equipment state.
- Simulation mode should be visibly separated from live operations.

Useful mental model

The visualisation is a window into the twin; it is not automatically the twin itself.

3. Conceptual Agentic Extension

What Does the Agent Add?


Good agent contribution

Interpret an operational goal, collect evidence, select bounded tools, compare options, ask for missing information, explain, and escalate.

Keep deterministic

Equipment protection, emergency stops, access enforcement, financial commitment, safety interlocks, and hard real-time control.

Agent Controller Loop Above the Twin



Loop engineering

The stopping rule, retry budget, evidence test, permission boundary, and escalation path matter as much as the prompt.

Example Typed Tool Surface

Tool	Input	Controlled result
get_asset_status	asset ID, time range	state and evidence references
get_zone_context	zone ID	occupancy, access, environment summary
dispatch_robot	mission, zone, priority	proposed/approved mission ID
create_ticket	asset, fault, evidence	auditable maintenance ticket
request_setpoint	zone, value, duration	approval request, not raw equipment write

Typed capability

A useful interface defines purpose, arguments, validation, output schema, side effects, permissions, and failure modes.

Robot Architecture: Separate Reasoning from Control

District agent

Creates a bounded mission from operational context.

Robot computer

Mission client, perception, mapping, navigation, video, ROS/ROS2.

Real-time controller

Motor control, local sensors, interlocks, emergency-stop behaviour.

Agent goal → validated mission → navigation/control → telemetry/evidence

Safety boundary

The LLM should not generate direct motor voltages or bypass local collision and emergency-stop logic.

State and Memory Need Time Semantics

- Current state: what is believed true now, with timestamp and source.
- Event history: what changed, when, and who or what changed it.
- Operational profile: stable structured facts about assets, zones, and policies.
- Recent summary: compressed context for the agent, never the sole source of truth.

Example

“Door was unlocked” is incomplete without the door ID, event time, validity interval, source, later superseding events, and confidence.

Ontology: Shared Meaning Across Systems

Working definition

An ontology defines domain concepts, identifiers, properties, relationships, and constraints so different systems can refer to the same things consistently.



- It may be represented in relational schemas, JSON schemas, knowledge graphs, RDF/OWL, or domain-specific models.
- The value is the agreed semantics, not one mandatory storage format.

4. Three Service Scenarios

Scenario 1: Controlled-Room Incident

Hypothetical event

A vision system reports a person in a controlled room. The application correlates the event with card access, room policy, time, and camera evidence.

- 1 retrieve structured detection and recent authorised access;
- 2 evaluate deterministic access/time-window rules;
- 3 ask for more evidence when events conflict;
- 4 optionally dispatch a robot through an approved mission service;
- 5 escalate to a human security operator with an evidence summary.

Do Not Make the LLM the Intrusion Detector

Specialised systems

Video analytics, badge systems, locks, and policy engines produce structured evidence and enforce access.


LLM/agent

Reconciles ambiguous evidence, explains the situation, drafts a warning, and recommends escalation within policy.

Critical action

Unlocking, detention, emergency response, or identity judgment requires explicit policy and appropriate human authority.

Scenario 2: Facility Fault and Private Manual



RAG value

Private manuals and maintenance histories may contain the needed evidence. Retrieval should expose the supporting section, not merely return an unsupported answer.

Facility Response Categories

Class	Response boundary	Example
A	robot may handle under fixed procedure	place a warning sign near a wet floor
B	robot may inspect; simple correction only if authorised	inspect a fallen sign or visible obstruction
C	qualified human inspection required	electrical circuit or lift mechanical fault
D	immediate emergency human response	fire, medical, or serious security event

Guardrail

The model may recommend a class; rules and policy validate the allowed action.

From Fault Diagnosis to Procurement

- The agent may retrieve evidence and draft a replacement-part recommendation.
- The workflow checks budget, vendor, approval authority, segregation of duties, and inventory.
- Purchase order, receiving, invoice match, and payment remain auditable transactions.



Boundary

The agent assists; the business workflow governs.

Scenario 3: Comfort, Cooling, and Energy

- Inputs: temperature, humidity, occupancy, timetable, event type, RSVP, weather, PV/storage/grid state, and policy.
- Predictive model: estimates load, attendance, or pre-cooling time.
- Agent: structures event context, compares scenarios, explains trade-offs, and requests approval for unusual operation.
- Building/energy controllers: enforce limits and execute safe setpoints.

Current PDD direction

JTC describes ODP-integrated district energy work including building systems, EV charging, district cooling, and a district-level smart grid.

Source: JTC, [District-level smart grid in PDD](#).

Simulation Before Live Control

Digital-twin experiment

Test a proposed schedule or setpoint using historical/current conditions; compare comfort, energy, cost, and risk.

Controlled deployment

Require validation and approval; constrain duration/range; observe; automatically revert or escalate if evidence is poor.

Why this matters

A twin can reduce risk, but simulation results are evidence—not proof that the physical system will behave identically.

Source: JTC, [Panasonic smart facility-management testbed in PDD, 2026](#).

5. Safety, Governance, and Reflection

Four Layers Against Agent Overreach

- 1 **Bind intent to identity and scope:** who requested what action, on which object, for what purpose?
- 2 **Least privilege and just-in-time access:** grant only the capability needed, for a short duration.
- 3 **Human confirmation for consequential actions:** present the exact target, effect, evidence, and rollback plan.
- 4 **Logging and anomaly detection:** preserve tool calls, approvals, state changes, and unexpected patterns.

Policy should be executable

Do not rely on a prompt saying “be careful.” Enforce permissions and validation outside the model.

Prevent One Agent from Damaging Another's Work

- Give each agent a separate workspace or branch; no shared unrestricted write access.
- Use ownership rules, path allowlists, file locks, and versioned artifacts.
- Require review/approval for deletion, overwrite, deployment, and irreversible operations.
- Prefer proposed patches or change sets to direct mutation of the source of truth.
- Record provenance: agent, model, prompt/task, input version, tool call, output, and reviewer.
- Run verification before merge; make rollback easy.


Coordination rule

Agents exchange tasks and artifacts through a controller or governed protocol; they should not “argue” by editing the same live files.

What Could Go Wrong?

Failure	Example	Engineering response
Stale state	agent acts on an old occupancy event	timestamps, validity, superseding events
Wrong identity	similar room/asset names	canonical IDs and ontology
Prompt injection	manual or ticket contains malicious text	treat retrieved content as data; policy outside LLM
Excess authority	diagnostic agent can alter live control	capability separation and temporary permission
Runaway loop	repeated robot dispatch or tickets	budgets, idempotency keys, stop/escalate rules
Silent conflict	two agents overwrite work	isolation, ownership, review, version control

Architecture Summary



Hybrid rule

Fixed systems detect and control; specialised AI perceives and predicts; agents coordinate and explain; workflows govern transactions; humans approve sensitive actions.

Employability: Skills Hidden in the Case Study

- systems integration: APIs, events, schemas, protocols, and adapters;
- agent engineering: tool interfaces, controller loops, evidence, stopping rules;
- data engineering: time-aware state, identity, provenance, and quality;
- cloud/backend: services, permissions, logging, observability, deployment;
- digital twin/graphics: spatial models, simulation, visualisation, interaction;
- robotics/embedded: ROS, perception, navigation, real-time control, safety;
- product/UX: human approval, operator trust, incident explanation, accessibility.

Reflection Questions

- 1 Which scenario genuinely needs an LLM, and which can be implemented as a deterministic workflow?
- 2 Where would you use direct APIs, CLI, events, MCP, ROS, or an industrial protocol?
- 3 What is the smallest safe tool surface for the agent?
- 4 What evidence must be fresh before dispatching a robot or changing a setpoint?
- 5 Which action requires human confirmation, and what exactly should the confirmation screen show?
- 6 How would you test the agent in simulation before allowing a limited live pilot?

Further Reading

- JTC: [Punggol Digital District overview](#).
- JTC: [Open Digital Platform—the digital backbone of PDD](#).
- JTC: [Five levels of the digital twin](#).
- JTC: [ODP for smart-city solutions](#).
- JTC: [District-level smart grid in PDD](#).
- JTC: [Smart facility-management testbed, 2026](#).
- IMDA: [Digital Connectivity Blueprint](#).

Final takeaway

The most valuable architecture is not “LLM controls the district.” It is a layered system in which the twin supplies trustworthy context and agents operate through narrow, governed, observable capabilities.